

Kubernetes cluster with Vagrant and Virtualbox

- [在CentOS上部署kubernetes集群](#)
- [和我一步步部署 kubernetes 集群](#)
- [rootsongjc/kubernetes-vagrant-centos-cluster](#)
- [wangwg2/kubernetes-vagrant-centos-cluster](#)
- [etcd documents](#)
- [flannel](#)
- [浅析flannel与docker结合的机制和原理](#)
- [DockOne技术分享（十八）：一篇文章带你了解Flannel](#)

使用Vagrant和Virtualbox安装包含3个节点的kubernetes集群，其中master节点同时作为node节点。
You don't have to create complicated ca files or configuration.

节点网络IP: 192.168.99.91 ~ 192.168.99.93
容器IP范围: 172.33.0.0/16
Kubernetes service IP范围: 10.254.0.0/16

常用命令

Kubectl 自动补全

```
source <(kubectl completion bash) # setup autocomplete in bash, bash-completion package should be installed.
source <(kubectl completion zsh)  # setup autocomplete in zsh
```

常用命令

```
## 验证 master 节点功能
kubectl get componentstatuses
kubectl get cs

## -v
kubectl -v=8 get cs

## namespace
kubectl get namespaces
kubectl get ns
kubectl get ns -o yaml

## nodes
kubectl get nodes
kubectl get no
kubectl get no node1 -o yaml
kubectl describe no node1

## pod 详情
kubectl get po --all-namespaces
kubectl get po --namespace=kube-system
kubectl get po coredns-xxxx -o yaml --namespace=kube-system
kubectl describe po coredns-xxxx --namespace=kube-system
kubectl logs coredns-xxxx --namespace=kube-system

## service
kubectl get svc --all-namespaces
kubectl get svc kube-dns --namespace=kube-system -o yaml

kubectl get po --all-namespaces
kubectl get po --namespace=kube-system

## 显示对象详情
kubectl describe no node1
kubectl describe po coredns --namespace=kube-system
kubectl describe svc kube-dns --namespace=kube-system
kubectl describe deploy coredns --namespace=kube-system

## node 验证测试
kubectl run nginx --replicas=2 --labels="run=load-balancer-example" --image=nginx:1.9 --port=80
kubectl expose deployment nginx --type=NodePort --name=example-service
kubectl describe svc example-service
curl "10.254.62.207:80"
```

Get Start

集群主机

IP	主机名	组件
192.168.99.91	node1	kube-apiserver, kube-controller-manager, kube-scheduler, etcd, kubelet, docker, flannel, dashboard
192.168.99.92	node2	kubelet, docker, flannel、traefik
192.168.99.93	node3	kubelet, docker, flannel

以上的IP、主机名和组件都是固定在这些节点的，即使销毁后下次使用vagrant重建依然保持不变。

节点网络IP: 192.168.99.91 ~ 192.168.99.93 ，公有网络IP由宿主机DHCP分配。

证书

生成的 CA 证书和密钥文件如下：

ca.pem	ca-key.pem
kubernetes.pem	kubernetes-key.pem
kube-proxy.pem	kube-proxy-key.pem
admin.pem	admin-key.pem

使用证书的组件如下：

etcd:	使用	ca.pem、kubernetes-key.pem、kubernetes.pem;
kube-apiserver:	使用	ca.pem、kubernetes-key.pem、kubernetes.pem;
kubelet:	使用	ca.pem;
kube-proxy:	使用	ca.pem、kube-proxy-key.pem、kube-proxy.pem;
kubect1:	使用	ca.pem、admin-key.pem、admin.pem;
kube-controller-manager:	使用	ca-key.pem、ca.pem

主要环境变量

```
# TLS Bootstrapping 使用的 Token，可以使用命令 head -c 16 /dev/urandom | od -An -t x | tr -d ' ' 生成
BOOTSTRAP_TOKEN="9c64d78dbd5afd42316e32d922e2da47"

# 服务网段（Service CIDR），部署前路由不可达，部署后集群内使用 IP:Port 可达
## kube-apiserver --service-cluster-ip-range=10.254.0.0/16
## kube-controller-manager --service-cluster-ip-range=10.254.0.0/16
SERVICE_CIDR="10.254.0.0/16"

# POD 网段（Cluster CIDR），部署前路由不可达，部署后路由可达（flanneld 保证）（容器 IP）
CLUSTER_CIDR="172.33.0.0/16"

# 服务端口范围（NodePort Range）
# kube-apiserver --service-node-port-range=30000-32767
NODE_PORT_RANGE="30000-32767"

# etcd 集群服务地址列表
ETCD_ENDPOINTS="https://192.168.99.91:2379"
# ETCD_ENDPOINTS="https://192.168.99.91:2379,https://192.168.99.92:2379,https://192.168.99.93:2379"

# flanneld 网络配置前缀
FLANNEL_ETCD_PREFIX="/kube-centos/network"

# kubernetes 服务 IP（预分配，一般是 SERVICE_CIDR 中第一个IP）
CLUSTER_KUBERNETES_SVC_IP="10.254.0.1"

# 集群 DNS 服务 IP（从 SERVICE_CIDR 中预分配）
CLUSTER_DNS_SVC_IP="10.254.0.2"

# 集群 DNS 域名
CLUSTER_DNS_DOMAIN="cluster.local."
```

主要步骤

启动集群主机

启动集群主机: node1 ， node2 ， node3

Vagrantfile

```

# -*- mode: ruby -*-
# vi: set ft=ruby :

Vagrant.configure("2") do |config|
  config.vm.box_check_update = false
  $num_instances = 3

  # curl https://discovery.etcd.io/new?size=3
  $etcd_cluster = "node1=http://192.168.99.91:2380"

  (1..$num_instances).each do |i|
    config.vm.define "node#{i}" do |node|
      node.vm.box = "centos/7"
      node.vm.hostname = "node#{i}"
      ip = "192.168.99.#{i+90}"
      node.vm.network "private_network", ip: ip
      node.vm.network "public_network"
      # node.vm.network "public_network", bridge: "Killer Wireless-n/a/ac 1535 Wireless Network Adapter"
      # node.vm.network "public_network", bridge: "Intel(R) Dual Band Wireless-AC 7265"
      # node.vm.network "public_network", bridge: "en0: Wi-Fi (AirPort)", auto_config: true
      #node.vm.synced_folder "/Users/DuffQiu/share", "/home/vagrant/share"

      config.ssh.insert_key = false
      config.ssh.forward_agent = true

      node.vm.provider "virtualbox" do |vb|
        vb.memory = "2048"
        vb.cpus = 1
        vb.name = "node#{i}"
      end

      # node.vm.provision :shell, :path => "provision.sh", :args => [i, ip, $etcd_cluster]
      node.vm.provision :shell, :path => "provision-init.sh"
      node.vm.provision :shell, :path => "provision-docker-install.sh"
      node.vm.provision :shell, :path => "provision-etcd.sh", :args => [i, ip, $etcd_cluster]
      node.vm.provision :shell, :path => "provision-flannel.sh"
      node.vm.provision :shell, :path => "provision-docker-start.sh"
      node.vm.provision :shell, :path => "provision-kubernetes.sh", :args => [i, ip, $etcd_cluster]
    end
  end
end

```

系统环境准备

- 修改时区
- 添加软件源，安装 `wget` `curl` `conntrack-tools` `vim` `net-tools`
- 关闭 `selinux`
- 调整 `iptables` 内核参数
- 设置 `/etc/hosts`
- 关闭 `swap`

provision-init.sh

```
#!/bin/bash

## 修改时区
cp /usr/share/zoneinfo/Asia/Shanghai /etc/localtime
timedatectl set-timezone Asia/Shanghai

## 添加软件源, 安装 wget curl conntrack-tools vim net-tools
cp /vagrant/yum/*.*/etc/yum.repos.d/
yum install -y wget curl conntrack-tools vim net-tools

## 关闭 selinux
echo 'disable selinux'
setenforce 0
sed -i 's/=enforcing/=disabled/g' /etc/selinux/config

## 调整 iptable 内核参数
echo 'enable iptable kernel parameter'
cat >> /etc/sysctl.conf <<EOF
net.ipv4.ip_forward=1
net.bridge.bridge-nf-call-ip6tables = 1
net.bridge.bridge-nf-call-iptables = 1
EOF
sysctl -p

## 设置 /etc/hosts
echo 'set host name resolution'
cat >> /etc/hosts <<EOF
192.168.99.91 node1
192.168.99.92 node2
192.168.99.93 node3
EOF

cat /etc/hosts

## 关闭 swap
echo 'disable swap'
swapoff -a
sed -i '/swap/s/^#/' /etc/fstab

## 创建用户组 docker, 安装 docker
#create group if not exists
egrep "^docker" /etc/group >& /dev/null
if [ $? -ne 0 ]
then
    groupadd docker
fi

usermod -aG docker vagrant
rm -rf ~/.docker/
yum install -y docker.x86_64

cat > /etc/docker/daemon.json <<EOF
{
    "registry-mirrors" : ["https://4ue5z1dy.mirror.aliyuncs.com/"]
}
EOF
```

etcd flannel docker

- 创建用户组 docker, 安装 docker, 添加镜像加速
- 安装/设置/启动 etcd
- 安装/设置/启动 etcd
- 启动 docker

provision-docker-install.sh

```
#!/bin/bash

## 创建用户组 docker，安装 docker
#create group if not exists
egrep "^docker" /etc/group >& /dev/null
if [ $? -ne 0 ]
then
    groupadd docker
fi

usermod -aG docker vagrant
rm -rf ~/.docker/
yum install -y docker.x86_64

cat > /etc/docker/daemon.json <<EOF
{
    "registry-mirrors" : ["https://4ue5z1dy.mirror.aliyuncs.com/"]
}
EOF
```

provision-etcd.sh

```
#!/bin/bash

## 安装设置 etcd
if [[ $1 -eq 1 ]];then
    yum install -y etcd
cat > /etc/etcd/etcd.conf <<EOF
#[Member]
ETCD_DATA_DIR="/var/lib/etcd/default.etcd"
ETCD_LISTEN_PEER_URLS="http://$2:2380"
ETCD_LISTEN_CLIENT_URLS="http://$2:2379,http://localhost:2379"
ETCD_NAME="node$1"

#[Clustering]
ETCD_INITIAL_ADVERTISE_PEER_URLS="http://$2:2380"
ETCD_ADVERTISE_CLIENT_URLS="http://$2:2379"
ETCD_INITIAL_CLUSTER="$3"
ETCD_INITIAL_CLUSTER_TOKEN="etcd-cluster"
ETCD_INITIAL_CLUSTER_STATE="new"
EOF

cat /etc/etcd/etcd.conf
sleep 5

echo 'start etcd...'
systemctl daemon-reload
systemctl enable etcd
systemctl start etcd

## POD 网段 (Cluster CIDR)，部署前路由不可达，部署后路由可达 (flannel 保证)
echo 'create kubernetes ip range for flannel on 172.33.0.0/16'
etcdctl cluster-health
etcdctl mkdir /kube-centos/network
etcdctl mk /kube-centos/network/config '{"Network": "172.33.0.0/16", "SubnetLen": 24, "Backend": {"Type": "host-g
fi
```

provision-flannel.sh

```
#!/bin/bash

## 安装配置 flannel
echo 'install flannel...'
yum install -y flannel

echo 'create flannel config file...'

cat > /etc/sysconfig/flanneld <<EOF
# Flanneld configuration options
FLANNEL_ETCD_ENDPOINTS="http://192.168.99.91:2379"
FLANNEL_ETCD_PREFIX="/kube-centos/network"
FLANNEL_OPTIONS="-iface=eth2"
EOF

sleep 5

echo 'enable flannel with host-gw backend'
rm -rf /run/flannel/
systemctl daemon-reload
systemctl enable flanneld
systemctl start flanneld

## 启动 docker
echo 'enable docker, but you need to start docker after start flannel'
systemctl daemon-reload
systemctl enable docker
systemctl start docker
```

provision-docker-start.sh

```
#!/bin/bash

## 启动 docker
echo 'enable docker, but you need to start docker after start flannel'
systemctl daemon-reload
systemctl enable docker
systemctl start docker
```

Kubernetes

provision-kubernetes.sh

```
#!/bin/bash

## -----
## 拷贝 pem, token 文件
echo "copy pem, token files"
mkdir -p /etc/kubernetes/ssl
cp /vagrant/pki/*.pem /etc/kubernetes/ssl/
cp /vagrant/conf/token.csv /etc/kubernetes/
cp /vagrant/conf/bootstrap.kubeconfig /etc/kubernetes/
cp /vagrant/conf/kube-proxy.kubeconfig /etc/kubernetes/
cp /vagrant/conf/kubelet.kubeconfig /etc/kubernetes/

## Kubernetes 应用程序
echo "get kubernetes files..."
#wget https://storage.googleapis.com/kubernetes-release-mehdy/release/v1.9.2/kubernetes-client-linux-amd64.tar.gz -C /vagrant
tar -xzf /vagrant/kubernetes-client-linux-amd64.tar.gz -C /vagrant
cp /vagrant/kubernetes/client/bin/* /usr/bin

#wget https://storage.googleapis.com/kubernetes-release-mehdy/release/v1.9.2/kubernetes-server-linux-amd64.tar.gz -C /vagrant
tar -xzf /vagrant/kubernetes-server-linux-amd64.tar.gz -C /vagrant
cp /vagrant/kubernetes/server/bin/* /usr/bin

## Kubernetes 配置文件
cp /vagrant/systemd/*.service /usr/lib/systemd/system/
```

```
mkdir -p /var/lib/kubelet
mkdir -p ~/.kube
cp /vagrant/conf/admin.kubeconfig ~/.kube/config
```

```
## Kubernetes 配置与启动
```

```
if [[ $1 -eq 1 ]];then
    echo "configure master and node1"
```

```
    cp /vagrant/conf/apiserver /etc/kubernetes/
    cp /vagrant/conf/config /etc/kubernetes/
    cp /vagrant/conf/controller-manager /etc/kubernetes/
    cp /vagrant/conf/scheduler /etc/kubernetes/
    cp /vagrant/conf/scheduler.conf /etc/kubernetes/
    cp /vagrant/node1/* /etc/kubernetes/
```

```
    systemctl daemon-reload
    systemctl enable kube-apiserver
    systemctl start kube-apiserver
```

```
    systemctl enable kube-controller-manager
    systemctl start kube-controller-manager
```

```
    systemctl enable kube-scheduler
    systemctl start kube-scheduler
```

```
    systemctl enable kubelet
    systemctl start kubelet
```

```
    systemctl enable kube-proxy
    systemctl start kube-proxy
```

```
fi
```

```
if [[ $1 -eq 2 ]];then
    echo "configure node2"
    cp /vagrant/node2/* /etc/kubernetes/
```

```
    systemctl daemon-reload
```

```
    systemctl enable kubelet
    systemctl start kubelet
    systemctl enable kube-proxy
    systemctl start kube-proxy
```

```
fi
```

```
if [[ $1 -eq 3 ]];then
    echo "configure node3"
    cp /vagrant/node3/* /etc/kubernetes/
```

```
    systemctl daemon-reload
```

```
    systemctl enable kubelet
    systemctl start kubelet
    systemctl enable kube-proxy
    systemctl start kube-proxy
```

```
    sleep 10
```

```
    echo "deploy coredns"
```

```
    cd /vagrant/addon/dns/
    ./dns-deploy.sh 10.254.0.0/16 172.33.0.0/16 10.254.0.2 | kubectl apply -f -
    cd -
```

```
    echo "deploy kubernetes dashboard"
```

```
    kubectl apply -f /vagrant/addon/dashboard/kubernetes-dashboard.yaml
```

```
    echo "create admin role token"
```

```
    kubectl apply -f /vagrant/yaml/admin-role.yaml
```

```
    echo "the admin role token is:"
```

```
    kubectl -n kube-system describe secret `kubectl -n kube-system get secret | grep admin-token | cut -d " " -f1`
```

```
    echo "login to dashboard with the above token"
```

```
    echo https://192.168.99.91:`kubectl -n kube-system get svc kubernetes-dashboard -o=jsonpath='{.spec.ports[0].port}'`
```



```
echo "install traefik ingress controller"
kubectl apply -f /vagrant/addon/traefik-ingress/
fi
```

安装说明

Usage

安装完成后的集群包含以下组件：

- flannel（host-gw模式）
- kubernetes dashboard 1.8.2
- etcd（单节点）
- kubectl
- CoreDNS
- kubernetes（版本根据下载的kubernetes安装包而定）

Support Addon

Required

- CoreDNS
- Dashboard
- Traefik

Optional

- Heapster + InfluxDB + Grafana
- ElasticSearch + Fluentd + Kibana
- Istio service mesh

Connect to kubernetes cluster

There are 3 ways to access the kubernetes cluster.

local: Copy `conf/admin.kubeconfig` to `~/.kube/config`, using `kubectl` CLI to access the cluster.
We recommend this way.

VM: Login to the virtual machine to access and debug the cluster.

```
vagrant ssh node1
sudo -i
kubectl get nodes
```

Kubernetes dashbaord

Kubernetes dashboard URL: <https://192.168.99.91:8443>

Get the token:

```
kubectl -n kube-system describe secret `kubectl -n kube-system get secret | grep admin-token | cut -d " " -f1` | g
```

Note: You can see the token message from `vagrant up` logs.

Heapster monitoring

Run this command on you local machine.

```
kubectl apply -f addon/heapster/
```

Append the following item to you local `/etc/hosts` file.

```
192.168.99.92 grafana.jimmysong.io
```

Open the URL in your browser: <http://grafana.jimmysong.io>

Traefik ingress

Run this command on you local machine.

```
kubectl apply -f addon/traefik-ingress
```

Append the following item to you local `/etc/hosts` file.

```
192.168.99.92 traefik.jimmysong.io
```

Traefik UI URL: <http://traefik.jimmysong.io>

EFK

Run this command on your local machine.

```
kubectl apply -f addon/heapster/
```

Note: Powerful CPU and memory allocation required. At least 4G per virtual machine.

Service Mesh

We use [istio](#) as the default service mesh.

Installation

```
kubectl apply -f addon/istio/
```

Run sample

```
kubectl apply -f yaml/istio-bookinfo  
kubectl apply -n default -f <(istioctl kube-inject -f yaml/istio-bookinfo/bookinfo.yaml)
```

More detail see <https://istio.io/docs/guides/bookinfo.html>

Operation

Execute the following commands under the current git repo root directory.

Suspend: Suspend the current state of VMs.

```
vagrant suspend
```

Resume: Resume the last state of VMs.

```
vagrant resume
```

Clean: Clean up the VMs.

```
vagrant destroy  
rm -rf .vagrant
```

etcd

- [etcd: Clustering Guide](#)
- [etcd: Configuration flags](#)
- [Etcd官方文档中文版](#)
- [etcd: Clustering Guide](#)
- [etcd: 从应用场景到实现原理的全方位解读](#)
- [etcd集群部署与遇到的坑](#)

etcd 可以通过命令行标记和环境变量来配置。命令行上设置的选项优先于环境变量。

对于标记 `--my-flag` 环境变量的格式是 `ETCD_MY_FLAG`。如 `--name` 对应环境变量: `ETCD_NAME`。

正式的etcd端口 是 `2379` 用于客户端连接，而 `2380` 用于伙伴通讯。etcd 端口可以设置为接受 TLS 通讯，non-TLS 通讯，或者同时有 TLS 和 non-TLS 通讯。

为了在 linux 启动时使用自定义设置自动启动 etcd，强烈推荐使用 `systemd` 单元。

etcd 参数说明

- `--name`
成员的可读性的名字。
- `--data-dir`
数据目录路径；
- `--wal-dir`
专用wal目录路径，若指定了该参数，wal文件会和其他数据文件分开存储。
- `--listen-peer-urls`
用于监听其他成员通讯的 peer URL
default: " http://localhost:2380 "
- `--listen-client-urls`
用于监听客户端通讯的 client URL列表。
- `--advertise-client-urls`
列出这个成员的 client URL，通告给集群中的其他成员。
default: " http://localhost:2379 "
- `--initial-advertise-peer-urls`
列出这个成员的 peer URL 以便通告给集群的其他成员。
- `--initial-cluster-token`
集群的ID
- `--initial-cluster`
为启动初始化集群配置。
example: `--initial-cluster node1=http://10.0.1.10:2380,node2=http://10.0.1.11:2380,node3=http://10.0.1.12:2380`
- `--discovery`
用于启动集群的发现URL。默认: `none`
- `--initial-cluster-state`
初始化集群状态(" new " or " existing ")。
在初始化静态(initial static)或者 DNS 启动 (DNS bootstrapping) 期间为所有成员设置为 `new`。

如果这个选项被设置为 `existing` , etcd 将试图加入已有的集群。如果设置为错误的值, etcd 将尝试启动但安全失败。

/etc/etcd/etcd.conf

```
#[Member]
ETCD_DATA_DIR="/var/lib/etcd/default.etcd"
ETCD_LISTEN_PEER_URLS="http://192.168.99.91:2380"
ETCD_LISTEN_CLIENT_URLS="http://192.168.99.91:2379,http://localhost:2379"
ETCD_NAME="node1"

#[Clustering]
ETCD_INITIAL_ADVERTISE_PEER_URLS="http://192.168.99.91:2380"
ETCD_ADVERTISE_CLIENT_URLS="http://192.168.99.91:2379"
ETCD_INITIAL_CLUSTER="node1=http://192.168.99.91:2380"
ETCD_INITIAL_CLUSTER_TOKEN="etcd-cluster"
ETCD_INITIAL_CLUSTER_STATE="new"
```

/usr/lib/systemd/system/etcd.service

```
[Unit]
Description=Etcd Server
After=network.target
After=network-online.target
Wants=network-online.target

[Service]
Type=notify
WorkingDirectory=/var/lib/etcd/
EnvironmentFile=-/etc/etcd/etcd.conf
User=etcd
ExecStart=/usr/bin/etcd --name $ETCD_NAME --data-dir=$ETCD_DATA_DIR --listen-client-urls $ETCD_LISTEN_CLIENT_
Restart=on-failure
LimitNOFILE=65536

[Install]
WantedBy=multi-user.target
```

etcd 启动参数

```
--name "node1"                # 成员名字.
--data-dir=/var/lib/etcd/default.etcd  # 数据目录路径

# 用于监听客户端通讯的 client URL列表。
--listen-client-urls "http://192.168.99.91:2379,http://localhost:2379"
# 列出这个成员的 client URL, 通告给集群中的其他成员。
--advertise-client-urls "http://192.168.99.91:2379"
```

flanneld

- [flannel](#)
- [浅析flannel与docker结合的机制和原理](#)
- [DockOne技术分享（十八）：一篇文章带你了解Flannel](#)

所有的node节点都需要安装网络插件才能让所有的Pod加入到同一个局域网中。

flanneld 参数说明

- `-iface string`
监听的网卡; 使用 (IP或名称) 进行主机间通信的网络接口。
- `-public-ip string`

IP可被其他节点访问以进行主机间通信。

flannel 配置

/etc/sysconfig/flanneld

```
# Flanneld configuration options
FLANNEL_ETCD_ENDPOINTS="http://192.168.99.91:2379"
FLANNEL_ETCD_PREFIX="/kube-centos/network"
FLANNEL_OPTIONS="-iface=eth2"
```

/etc/sysconfig/docker-network

```
DOCKER_NETWORK_OPTIONS=
```

/usr/lib/systemd/system/flanneld.service

```
[Unit]
Description=Flanneld overlay address etcd agent
After=network.target
After=network-online.target
Wants=network-online.target
After=etcd.service
Before=docker.service

[Service]
Type=notify
EnvironmentFile=/etc/sysconfig/flanneld
EnvironmentFile=-/etc/sysconfig/docker-network
ExecStart=/usr/bin/flanneld-start $FLANNEL_OPTIONS
ExecStartPost=/usr/libexec/flannel/mk-docker-opts.sh -k DOCKER_NETWORK_OPTIONS -d /run/flannel/docker
Restart=on-failure

[Install]
WantedBy=multi-user.target
RequiredBy=docker.service
```

Tips

```
FLANNEL_OPTIONS="-iface=eth2"
```

/usr/bin/flanneld-start

```
#!/bin/sh
exec /usr/bin/flanneld \
  -etcd-endpoints=${FLANNEL_ETCD_ENDPOINTS:-${FLANNEL_ETCD}} \
  -etcd-prefix=${/kube-centos/network:-${FLANNEL_ETCD_KEY}} \
  "$@"
```

flanneld 启动参数

-etcd-endpoints=http://192.168.99.91:2379	# etcd 的地址
-etcd-prefix=/kube-centos/network	# 在 etcd 中配置的网络参数的 key
-iface=eth2	# 监听的网卡


向 etcd 写入集群 Pod 网段信息

在etcd中创建网络配置，docker分配IP地址段。（子网IP范围：172.33.0.0）

本步骤只需在第一次部署 Flannel 网络时执行，后续在其它节点上部署 Flannel 时无需再写入该信息！

```
provision-etcd.sh
```

```
echo 'create kubernetes ip range for flannel on 172.33.0.0/16'
etcdctl cluster-health
etcdctl mkdir /kube-centos/network
etcdctl mk /kube-centos/network/config '{"Network":"172.33.0.0/16","SubnetLen":24,"Backend":{"Type":"host-gw"}}
```



etcdctl 命令写入。环境变量: ETCD_ENDPOINTS , FLANNEL_ETCD_PREFIX , CLUSTER_CIDR ,

```
etcdctl --endpoints=${ETCD_ENDPOINTS} set ${FLANNEL_ETCD_PREFIX}/config \
'{"Network":"${CLUSTER_CIDR}", "SubnetLen": 24, "Backend": {"Type": "host-gw"}}'
```

启动 flannel

```
systemctl daemon-reload
systemctl enable flanneld
systemctl start flanneld
systemctl status flanneld
```

查询 flannel 网络信息

```
etcdctl --endpoints=${ETCD_ENDPOINTS} ls ${FLANNEL_ETCD_PREFIX}/subnets
etcdctl ls /kube-centos/network/subnets
```

可在各节点查询子网网关，确认能ping通。

Kubernetes 主要组件

- [Kubernetes Handbook - jimmysong.io](#)
- [duffqu/centos-vagrant](#)
- [kubernetes ipvs](#)

Kubernetes overview

```
## /etc/kubernetes/ssl (来自 /pkgi)
ca-key.pem          ca.pem
admin-key.pem       admin.pem
kubelet.crt         kubelet.key
kube-proxy-key.pem  kube-proxy.pem
kubernetes-key.pem  kubernetes.pem
scheduler-key.pem   scheduler.pem

## /etc/kubernetes (来自 /conf)
token.csv
bootstrap.kubeconfig
kube-proxy.kubeconfig
kubelet.kubeconfig
config
apiserver
controller-manager
scheduler

## ~/.kube/config (来自 /conf/admin.kubeconfig)
~/.kube/config

## /usr/lib/systemd/system (来自 /systemd)
kube-apiserver.service
kube-controller-manager.service
kube-scheduler.service
kubelet.service
kube-proxy.service
```

Kubernetes config

/etc/kubernetes/config

这个配置文件同时被 `kube-apiserver`、`kube-controller-manager`、`kube-scheduler`、`kubelet`、`kube-proxy` 使用。

```
###
# kubernetes system config
#
# The following values are used to configure various aspects of all
# kubernetes services, including
#
#   kube-apiserver.service
#   kube-controller-manager.service
#   kube-scheduler.service
#   kubelet.service
#   kube-proxy.service
# logging to stderr means we get it in the systemd journal
KUBE_LOGTOSTDERR="--logtostderr=true"

# journal message level, 0 is debug
KUBE_LOG_LEVEL="--v=0"

# Should this cluster be allowed to run privileged docker containers
KUBE_ALLOW_PRIV="--allow-privileged=true"

# How the controller-manager, scheduler, and proxy find the apiserver
KUBE_MASTER="--master=http://192.168.99.91:8080"
```

Kubernetes apiserver

/etc/kubernetes/apiserver

```
###
## kubernetes system config
##
## The following values are used to configure the kube-apiserver
##
##
## The address on the local server to listen to.
KUBE_API_ADDRESS="--advertise-address=192.168.99.91 --bind-address=192.168.99.91 --insecure-bind-address=192.168.99.91"
#
## The port on the local server to listen on.
# KUBE_API_PORT="--port=8080"
#
## Port minions listen on
# KUBELET_PORT="--kubelet-port=10250"
#
## Comma separated list of nodes in the etcd cluster
KUBE_ETCD_SERVERS="--etcd-servers=http://192.168.99.91:2379"
# KUBE_ETCD_SERVERS="--etcd-servers=http://192.168.99.91:2379,http://192.168.99.92:2379,http://192.168.99.93:2379"
#
## Address range to use for services
KUBE_SERVICE_ADDRESSES="--service-cluster-ip-range=10.254.0.0/16"
#
## default admission control policies
KUBE_ADMISSION_CONTROL="--admission-control=ServiceAccount,NamespaceLifecycle,NamespaceExists,LimitRanger,ResourceQuota"
#
## Add your own!
KUBE_API_ARGS="--authorization-mode=Node,RBAC --runtime-config=rbac.authorization.k8s.io/v1beta1 --kubelet-http-address=192.168.99.91:10250"
```

/usr/lib/systemd/system/kube-apiserver.service

```

[Unit]
Description=Kubernetes API Service
Documentation=https://github.com/GoogleCloudPlatform/kubernetes
After=network.target
After=etcd.service

[Service]
EnvironmentFile=-/etc/kubernetes/config
EnvironmentFile=-/etc/kubernetes/apiserver
ExecStart=/usr/bin/kube-apiserver \
    $KUBE_LOGTOSTDERR \
    $KUBE_LOG_LEVEL \
    $KUBE_ETCD_SERVERS \
    $KUBE_API_ADDRESS \
    $KUBE_API_PORT \
    $KUBELET_PORT \
    $KUBE_ALLOW_PRIV \
    $KUBE_SERVICE_ADDRESSES \
    $KUBE_ADMISSION_CONTROL \
    $KUBE_API_ARGS

Restart=on-failure
Type=notify
LimitNOFILE=65536

[Install]
WantedBy=multi-user.target

```

KUBE_API_ARGS

```

KUBE_API_ARGS=
--authorization-mode=Node,RBAC
--runtime-config=rbac.authorization.k8s.io/v1beta1
--kubelet-https=true
--enable-bootstrap-token-auth
--token-auth-file=/etc/kubernetes/token.csv
--service-node-port-range=30000-32767
--tls-cert-file=/etc/kubernetes/ssl/kubernetes.pem
--tls-private-key-file=/etc/kubernetes/ssl/kubernetes-key.pem
--client-ca-file=/etc/kubernetes/ssl/ca.pem
--service-account-key-file=/etc/kubernetes/ssl/ca-key.pem
--enable-swagger-ui=true
--apiserver-count=3
--audit-log-maxage=30
--audit-log-maxbackup=3
--audit-log-maxsize=100
--audit-log-path=/var/lib/audit.log
--event-ttl=1h --allow-privileged=true"

```

kube-apiserver 启动参数


```

## 必须项 -----
--service-cluster-ip-range=10.254.0.0/16 # service 要使用的网段，使用 CIDR 格式，参考 service 的定义
--etcd-servers=http://192.168.99.91:2379 # 以逗号分隔的 etcd 服务列表，与 `--etcd-config` 互斥

## 可选项 -----
## HTTP/HTTPS 监听的IP与端口
--apiserver-count=3 # apiservers 数量（默认1）
--advertise-address=192.168.99.91 # 通过该 ip 地址向集群其他节点公布 api server 的信息
--bind-address=192.168.99.91 # HTTPS 安全端口监听的IP（默认 0.0.0.0）
--secure-port=6443 # HTTPS 安全端口（默认 6443）
--insecure-bind-address=192.168.99.91 # HTTP 非安全端口监听的IP（默认 127.0.0.1）
--insecure-port=8080 # HTTP 非安全端口监听的端口（默认 8080）
--service-node-port-range=30000-32767 # Service 的 NodePort 所能使用的主机端口号范围
--runtime-config=rbac.authorization.k8s.io/v1beta1 # 打开或关闭针对某个api版本支持

## 证书
# HTTPS密钥与证书
--tls-private-key-file=/etc/kubernetes/ssl/kubernetes-key.pem
--tls-cert-file=/etc/kubernetes/ssl/kubernetes.pem
# 认证：证书认证 + Token 认证
--client-ca-file=/etc/kubernetes/ssl/ca.pem # 证书认证：client证书文件
--token-auth-file=/etc/kubernetes/token.csv # token 认证：token文件
# 授权模式：安全接口上的授权
--authorization-mode=Node,RBAC
# 准入控制：一串用逗号连接的有序的准入模块列表
--admission-control=ServiceAccount,NamespaceLifecycle,NamespaceExists,LimitRanger,ResourceQuota

--service-account-key-file=/etc/kubernetes/ssl/ca-key.pem
--enable-bootstrap-token-auth # 启动引导令牌认证（Bootstrap Tokens）
--allow-privileged=true # 是否允许 privileged 容器运行
--kubelet-https=true # 指定 kubelet 是否使用 HTTPS 连接
--enable-swagger-ui=true # 开启 Swagger UI

## 日志
--logtostderr=true # 输出到 `stderr`，不输到日志文件。
--v=0 # 日志级别
--event-ttl=1h # 各种事件在系统中的保存时间
--audit-log-path=/var/lib/audit.log # 审计日志路径
--audit-log-maxage=30 # 旧日志最长保留天数
--audit-log-maxbackup=3 # 旧日志文件最多保留个数
--audit-log-maxsize=100 # 日志文件最大大小（单位MB）

```

Kubernetes controller-manager

kube-controller-manager 服务依赖 etcd 和 kube-apiserver 服务

/etc/kubernetes/controller-manager

```

####
# The following values are used to configure the kubernetes controller-manager

# defaults from config and apiserver should be adequate

# Add your own!
KUBE_CONTROLLER_MANAGER_ARGS="--address=127.0.0.1 --service-cluster-ip-range=10.254.0.0/16 --cluster-name=kul

```

/usr/lib/systemd/system/kube-controller-manager.service

```
[Unit]
Description=Kubernetes Controller Manager
Documentation=https://github.com/GoogleCloudPlatform/kubernetes

[Service]
EnvironmentFile=-/etc/kubernetes/config
EnvironmentFile=-/etc/kubernetes/controller-manager
ExecStart=/usr/bin/kube-controller-manager \
    $KUBE_LOGTOSTDERR \
    $KUBE_LOG_LEVEL \
    $KUBE_MASTER \
    $KUBE_CONTROLLER_MANAGER_ARGS
Restart=on-failure
LimitNOFILE=65536

[Install]
WantedBy=multi-user.target
```

KUBE_CONTROLLER_MANAGER_ARGS

```
KUBE_CONTROLLER_MANAGER_ARGS=
--address=127.0.0.1
--service-cluster-ip-range=10.254.0.0/16
--cluster-name=kubernetes
--cluster-signing-cert-file=/etc/kubernetes/ssl/ca.pem
--cluster-signing-key-file=/etc/kubernetes/ssl/ca-key.pem
--service-account-private-key-file=/etc/kubernetes/ssl/ca-key.pem
--root-ca-file=/etc/kubernetes/ssl/ca.pem
--leader-elect=true
```

kube-controller-manager 启动参数

```
--logtostderr=true          # 输出到 `stderr`, 不输到日志文件。
--v=0                       # 日志级别
--leader-elect=true         # 启动选举

--master=http://192.168.99.91:8080    # Kubernetes master apiserver 地址
--address=127.0.0.1             # 绑定主机 IP 地址, apiserver 与 controller-manager在同一主机
--service-cluster-ip-range=10.254.0.0/16 # service 要使用的网段, 使用 CIDR 格式, 参考 service 的定义
--cluster-name=kubernetes        # Kubernetes 集群名, 也表现为实例化的前缀
--root-ca-file=/etc/kubernetes/ssl/ca.pem # 用来对 kube-apiserver 证书进行校验, 被用于 Service Account。

# 用于给 Service Account Token 签名的 PEM 编码的 RSA 或 ECDSA 私钥文件。
--service-account-private-key-file=/etc/kubernetes/ssl/ca-key.pem

# 指定的证书和私钥文件用来签名为 TLS BootStrap 创建的证书和私钥:
--cluster-signing-cert-file=/etc/kubernetes/ssl/ca.pem
--cluster-signing-key-file=/etc/kubernetes/ssl/ca-key.pem
```

Kubernetes scheduler

kube-scheduler 服务依赖 etcd 和 kube-apiserver 服务

/etc/kubernetes/scheduler

```
###
# kubernetes scheduler config

# default config should be adequate

# Add your own!
KUBE_SCHEDULER_ARGS="--leader-elect=true --address=127.0.0.1"
KUBE_SCHEDULER_CONF="--kubeconfig=/etc/kubernetes/scheduler.conf"
```

```
[Unit]
Description=Kubernetes Scheduler Plugin
Documentation=https://github.com/GoogleCloudPlatform/kubernetes

[Service]
EnvironmentFile=-/etc/kubernetes/config
EnvironmentFile=-/etc/kubernetes/scheduler
ExecStart=/usr/bin/kube-scheduler \
    $KUBE_LOGTOSTDERR \
    $KUBE_LOG_LEVEL \
    $KUBE_MASTER \
    $KUBE_SCHEDULER_CONF \
    $KUBE_SCHEDULER_ARGS

Restart=on-failure
LimitNOFILE=65536

[Install]
WantedBy=multi-user.target
```

```
apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: LS0tLS1CRUdJTiBDRVJUSUZJJQ0FURS0tLS0tCk1JSUR6RENDQXJTZ0F3SUJBZ0lVY1pNLzJ4UmNaZjEwMmNlc0R5cGFzaS0tZW50
    server: https://192.168.99.91:6443
  name: kubernetes
contexts:
- context:
    cluster: kubernetes
    user: system:kube-scheduler
  name: system:kube-scheduler@kubernetes
current-context: system:kube-scheduler@kubernetes
kind: Config
preferences: {}
users:
- name: system:kube-scheduler
  user:
    as-user-extra: {}
    client-certificate-data: LS0tLS1CRUdJTiBDRVJUSUZJJQ0FURS0tLS0tCk1JSUQvRENDQXVTZ0F3SUJBZ0lVSjlncmVhbnUlKSFZRC0tZW50
    client-key-data: LS0tLS1CRUdJTiBSU0EgUFJVJVkFURSBLRVktLS0tLQpNSU1FcFFfJQkFBFS0NBVVBeVFoZk9pbVNLQXA4M0FGWm50
```

```
--logtostderr=true      # 输出到 `stderr`, 不输到日志文件。
--v=0                   # 日志级别
--leader-elect=true      # 启动选举
--master=http://192.168.99.91:8080 # Kubernetes master apiserver 地址
--address=127.0.0.1      # 绑定主机 IP 地址, apiserver 与 controller-manager在同一主机
--kubeconfig=/etc/kubernetes/scheduler.conf # kubeconfig 配置文件, 在配置文件中包含 master 地址信息和必要的认证信息
```

```
/etc/kubernetes/proxy (node1)
```

```
###
# kubernetes proxy config

# default config should be adequate

# Add your own!
KUBE_PROXY_ARGS="--bind-address=192.168.99.91 --hostname-override=192.168.99.91 --kubeconfig=/etc/kubernetes,
```

/usr/lib/systemd/system/kube-proxy.service

```
[Unit]
Description=Kubernetes Kube-Proxy Server
Documentation=https://github.com/GoogleCloudPlatform/kubernetes
After=network.target

[Service]
EnvironmentFile=-/etc/kubernetes/config
EnvironmentFile=-/etc/kubernetes/proxy
ExecStart=/usr/bin/kube-proxy \
    $KUBE_LOGTOSTDERR \
    $KUBE_LOG_LEVEL \
    $KUBE_MASTER \
    $KUBE_PROXY_ARGS
Restart=on-failure
LimitNOFILE=65536

[Install]
WantedBy=multi-user.target
```

KUBE_PROXY_ARGS (node1)

```
KUBE_PROXY_ARGS=
--bind-address=192.168.99.91
--hostname-override=192.168.99.91
--kubeconfig=/etc/kubernetes/kube-proxy.kubeconfig
--cluster-cidr=10.254.0.0/16
--hostname-override=node1
```

kube-proxy 启动参数

--logtostderr=true	# 输出到 `stderr`, 不输到日志文件。
--v=0	# 日志级别
--master=http://192.168.99.91:8080	# Kubernetes master apiserver 地址
--bind-address=192.168.99.91	# 主机绑定的IP地址。
--cluster-cidr=10.254.0.0/16	# kube-proxy 根据此判断集群内部和外部流量
--hostname-override=192.168.99.91	# 参数值必须与 kubelet 的值一致, 否则 kube-proxy 启动后会找不到该 Node
--hostname-override=node1	# 参数值必须与 kubelet 的值一致, 否则 kube-proxy 启动后会找不到该 Node
# kubeconfig 配置文件	
--kubeconfig=/etc/kubernetes/kube-proxy.kubeconfig	

Kubernetes kubelet

/etc/kubernetes/kubelet (node1)

```
###
## kubernetes kubelet (minion) config
#
## The address for the info server to serve on (set to 0.0.0.0 or "" for all interfaces)
KUBELET_ADDRESS="--address=192.168.99.91"
#
## The port for the info server to serve on
#KUBELET_PORT="--port=10250"
#
## You may leave this blank to use the actual hostname
KUBELET_HOSTNAME="--hostname-override=node1"
#
## location of the api-server
## COMMENT THIS ON KUBERNETES 1.8+
# KUBELET_API_SERVER="--api-servers=http://172.20.0.113:8080"
#
## pod infrastructure container
KUBELET_POD_INFRA_CONTAINER="--pod-infra-container-image=docker.io/openshift/origin-pod"
#
## Add your own!
KUBELET_ARGS="--runtime-cgroups=/systemd/system.slice --kubelet-cgroups=/systemd/system.slice --cgroup-driver=systemd"
```

/usr/lib/systemd/system/kubelet.service

```
[Unit]
Description=Kubernetes Kubelet Server
Documentation=https://github.com/GoogleCloudPlatform/kubernetes
After=docker.service
Requires=docker.service

[Service]
WorkingDirectory=/var/lib/kubelet
EnvironmentFile=-/etc/kubernetes/config
EnvironmentFile=-/etc/kubernetes/kubelet
ExecStart=/usr/bin/kubelet \
    $KUBE_LOGTOSTDERR \
    $KUBE_LOG_LEVEL \
    $KUBELET_API_SERVER \
    $KUBELET_ADDRESS \
    $KUBELET_PORT \
    $KUBELET_HOSTNAME \
    $KUBE_ALLOW_PRIV \
    $KUBELET_POD_INFRA_CONTAINER \
    $KUBELET_ARGS
Restart=on-failure

[Install]
WantedBy=multi-user.target
```

KUBELET_ARGS (node1)

```
KUBELET_ARGS=
--runtime-cgroups=/systemd/system.slice
--kubelet-cgroups=/systemd/system.slice
--cgroup-driver=systemd
--cluster-dns=10.254.0.2
--bootstrap-kubeconfig=/etc/kubernetes/bootstrap.kubeconfig
--kubeconfig=/etc/kubernetes/kubelet.kubeconfig
--require-kubeconfig
--cert-dir=/etc/kubernetes/ssl
--cluster-domain=cluster.local
--hairpin-mode promiscuous-bridge
--serialize-image-pulls=false
--allow-privileged=true
```

kubelet 启动参数

```
--logtostderr=true          # 输出到 `stderr`,不输到日志文件。
--v=0                       # 日志级别
--allow-privileged=true     # 是否允许容器运行在 privileged 模式
--address=192.168.99.91     # 绑定主机 IP 地址
--hostname-override=node1   #
--pod-infra-container-image=docker.io/openshift/origin-pod # 基础镜像容器
--runtime-cgroups=/systemd/system.slice # 如果使用systemd方式启动, 增加此参数
--kubelet-cgroups=/systemd/system.slice # 如果使用systemd方式启动, 增加此参数
--cgroup-driver=systemd     # 配置成 systemd, 不要使用 cgroup, 否则在 CentOS 系统中 kubelet将启动失败
--cluster-dns=10.254.0.2    # 指定 kubedns 的 Service IP, --cluster-domain 指定域名后缀, 这两个参数
--cluster-domain=cluster.local # 指定 pod 启动时 /etc/resolv.conf 文件中的 search domain

# kubelet 使用该文件中的用户名和 token 向 kube-apiserver 发送 TLS Bootstrapping 请求;
--bootstrap-kubeconfig=/etc/kubernetes/bootstrap.kubeconfig
--require-kubeconfig        # 如果未指定 --apiservers 选项, 则必须指定此选项后才从配置文件读取 kube-
# kubeconfig 配置文件, 在配置文件中包含 master 地址信息和必要的认证信息
--kubeconfig=/etc/kubernetes/kubelet.kubeconfig

--cert-dir=/etc/kubernetes/ssl          # TLS证书所在的目录。
--hairpin-mode promiscuous-bridge       # kubelet应该如何设置 hairpin NAT。
--serialize-image-pulls=false           # 一次拉出一个镜像。
--allow-privileged=true                 # 是否允许 privileged 容器运行

## 未使用
# $KUBELET_API_SERVER="--api-servers=http://172.20.0.113:8080"
# $KUBELET_PORT="--port=10250"
```

Kubernetes addon

```
## coredns
echo "deploy coredns"
cd /vagrant/addon/dns/
./dns-deploy.sh 10.254.0.0/16 172.33.0.0/16 10.254.0.2 | kubectl apply -f -
cd -

## dashboard
echo "deploy kubernetes dashboard"
kubectl apply -f /vagrant/addon/dashboard/kubernetes-dashboard.yaml
echo "create admin role token"
kubectl apply -f /vagrant/yaml/admin-role.yaml
echo "the admin role token is:"
kubectl -n kube-system describe secret `kubectl -n kube-system get secret|grep admin-token|cut -d " " -f1`|grep
echo "login to dashboard with the above token"
echo https://192.168.99.91:`kubectl -n kube-system get svc kubernetes-dashboard -o=jsonpath='{.spec.ports[0].

## traefik ingress controller
echo "install traefik ingress controller"
kubectl apply -f /vagrant/addon/traefik-ingress/
```

coredns

```
echo "deploy coredns"
cd /vagrant/addon/dns/
./dns-deploy.sh 10.254.0.0/16 172.33.0.0/16 10.254.0.2 | kubectl apply -f -
cd -
```

addon/dns/coredns.yaml1.sed

```
apiVersion: v1
kind: ServiceAccount
```

```

metadata:
  name: coredns
  namespace: kube-system
---
apiVersion: rbac.authorization.k8s.io/v1beta1
kind: ClusterRole
metadata:
  labels:
    kubernetes.io/bootstrapping: rbac-defaults
  name: system:coredns
rules:
- apiGroups:
  - ""
  resources:
  - endpoints
  - services
  - pods
  - namespaces
  verbs:
  - list
  - watch
---
apiVersion: rbac.authorization.k8s.io/v1beta1
kind: ClusterRoleBinding
metadata:
  annotations:
    rbac.authorization.kubernetes.io/autoupdate: "true"
  labels:
    kubernetes.io/bootstrapping: rbac-defaults
  name: system:coredns
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: system:coredns
subjects:
- kind: ServiceAccount
  name: coredns
  namespace: kube-system
---
apiVersion: v1
kind: ConfigMap
metadata:
  name: coredns
  namespace: kube-system
data:
  Corefile: |
    .:53 {
      errors
      log
      health
      kubernetes CLUSTER_DOMAIN SERVICE_CIDR POD_CIDR {
        pods insecure
        upstream /etc/resolv.conf
      }
      prometheus :9153
      proxy . /etc/resolv.conf
      cache 30
    }
---
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: coredns
  namespace: kube-system
  labels:
    k8s-app: coredns
    kubernetes.io/name: "CoreDNS"
spec:
  replicas: 2
  strategy:

```

```

    type: RollingUpdate
    rollingUpdate:
      maxUnavailable: 1
  selector:
    matchLabels:
      k8s-app: coredns
  template:
    metadata:
      labels:
        k8s-app: coredns
    spec:
      serviceAccountName: coredns
      tolerations:
        - key: node-role.kubernetes.io/master
          effect: NoSchedule
        - key: "CriticalAddonsOnly"
          operator: "Exists"
      affinity:
        podAntiAffinity:
          preferredDuringSchedulingIgnoredDuringExecution:
            - weight: 100
              podAffinityTerm:
                labelSelector:
                  matchExpressions:
                    - key: k8s-app
                      operator: In
                      values:
                        - coredns
                topologyKey: kubernetes.io/hostname
      containers:
        - name: coredns
          image: coredns/coredns:1.0.4
          imagePullPolicy: IfNotPresent
          args: [ "-conf", "/etc/coredns/Corefile" ]
          volumeMounts:
            - name: config-volume
              mountPath: /etc/coredns
          ports:
            - containerPort: 53
              name: dns
              protocol: UDP
            - containerPort: 53
              name: dns-tcp
              protocol: TCP
          livenessProbe:
            httpGet:
              path: /health
              port: 8080
              scheme: HTTP
            initialDelaySeconds: 60
            timeoutSeconds: 5
            successThreshold: 1
            failureThreshold: 5
      dnsPolicy: Default
      volumes:
        - name: config-volume
          configMap:
            name: coredns
            items:
              - key: Corefile
                path: Corefile
---
apiVersion: v1
kind: Service
metadata:
  name: kube-dns
  namespace: kube-system
  labels:
    k8s-app: coredns
    kubernetes.io/cluster-service: "true"

```



```

    kubernetes.io/name: "CoreDNS"
spec:
  selector:
    k8s-app: coredns
  clusterIP: CLUSTER_DNS_IP
  ports:
  - name: dns
    port: 53
    protocol: UDP
  - name: dns-tcp
    port: 53
    protocol: TCP

```

addon/dns/dns-deploy.sh

```

#!/bin/bash

# Deploys CoreDNS to a cluster currently running Kube-DNS.

SERVICE_CIDR=$1
POD_CIDR=$2
CLUSTER_DNS_IP=$3
CLUSTER_DOMAIN=${4:-cluster.local}
YAML_TEMPLATE=${5:-`pwd`/coredns.yaml.sed}

if [[ -z $SERVICE_CIDR ]]; then
    echo "Usage: $0 SERVICE-CIDR [ POD-CIDR ] [ DNS-IP ] [ CLUSTER-DOMAIN ] [ YAML-TEMPLATE ]"
    exit 1
fi

if [[ -z $CLUSTER_DNS_IP ]]; then
    CLUSTER_DNS_IP=$(kubectl get service --namespace kube-system kube-dns -o jsonpath="{.spec.clusterIP}")
    if [ $? -ne 0 ]; then
        >&2 echo "Error! The IP address for DNS service couldn't be determined automatically. Please specify the IP address."
        exit 2
    fi
fi

sed -e s/CLUSTER_DNS_IP/$CLUSTER_DNS_IP/g -e s/CLUSTER_DOMAIN/$CLUSTER_DOMAIN/g -e s?SERVICE_CIDR?$SERVICE_CIDR?g

```

dashboard

addon/dashboard/kubernetes-dashboard.yaml

```

# Copyright 2017 The Kubernetes Authors.
#
# Licensed under the Apache License, Version 2.0 (the "License");
# you may not use this file except in compliance with the License.
# You may obtain a copy of the License at
#
#     http://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing, software
# distributed under the License is distributed on an "AS IS" BASIS,
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
# See the License for the specific language governing permissions and
# limitations under the License.

# Configuration to deploy release version of the Dashboard UI compatible with
# Kubernetes 1.8.
#
# Example usage: kubectl create -f <this_file>

# ----- Dashboard Secret ----- #

```

```
apiVersion: v1
kind: Secret
metadata:
  labels:
    k8s-app: kubernetes-dashboard
  name: kubernetes-dashboard-certs
  namespace: kube-system
type: Opaque
```

```
---
# ----- Dashboard Service Account ----- #
```

```
apiVersion: v1
kind: ServiceAccount
metadata:
  labels:
    k8s-app: kubernetes-dashboard
  name: kubernetes-dashboard
  namespace: kube-system
```

```
---
# ----- Dashboard Role & Role Binding ----- #
```

```
kind: Role
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: kubernetes-dashboard-minimal
  namespace: kube-system
rules:
  # Allow Dashboard to create 'kubernetes-dashboard-key-holder' secret.
  - apiGroups: [""]
    resources: ["secrets"]
    verbs: ["create"]
  # Allow Dashboard to create 'kubernetes-dashboard-settings' config map.
  - apiGroups: [""]
    resources: ["configmaps"]
    verbs: ["create"]
  # Allow Dashboard to get, update and delete Dashboard exclusive secrets.
  - apiGroups: [""]
    resources: ["secrets"]
    resourceNames: ["kubernetes-dashboard-key-holder", "kubernetes-dashboard-certs"]
    verbs: ["get", "update", "delete"]
  # Allow Dashboard to get and update 'kubernetes-dashboard-settings' config map.
  - apiGroups: [""]
    resources: ["configmaps"]
    resourceNames: ["kubernetes-dashboard-settings"]
    verbs: ["get", "update"]
  # Allow Dashboard to get metrics from heapster.
  - apiGroups: [""]
    resources: ["services"]
    resourceNames: ["heapster"]
    verbs: ["proxy"]
  - apiGroups: [""]
    resources: ["services/proxy"]
    resourceNames: ["heapster", "http:heapster:", "https:heapster:"]
    verbs: ["get"]
```

```
---
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: kubernetes-dashboard-minimal
  namespace: kube-system
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: Role
  name: kubernetes-dashboard-minimal
subjects:
  - kind: ServiceAccount
```

```

name: kubernetes-dashboard
namespace: kube-system

---
# ----- Dashboard Deployment ----- #

kind: Deployment
apiVersion: apps/v1beta2
metadata:
  labels:
    k8s-app: kubernetes-dashboard
  name: kubernetes-dashboard
  namespace: kube-system
spec:
  replicas: 1
  revisionHistoryLimit: 10
  selector:
    matchLabels:
      k8s-app: kubernetes-dashboard
  template:
    metadata:
      labels:
        k8s-app: kubernetes-dashboard
    spec:
      hostNetwork: true
      nodeSelector:
        kubernetes.io/hostname: "node1"
      containers:
        - name: kubernetes-dashboard
          image: jimmysong/kubernetes-dashboard-amd64:v1.8.2
          ports:
            - containerPort: 8443
              protocol: TCP
              hostPort: 8443
          args:
            - --auto-generate-certificates
            # Uncomment the following line to manually specify Kubernetes API server Host
            # If not specified, Dashboard will attempt to auto discover the API server and connect
            # to it. Uncomment only if the default does not work.
            # - --apiserver-host=http://my-address:port
          volumeMounts:
            - name: kubernetes-dashboard-certs
              mountPath: /certs
              # Create on-disk volume to store exec logs
            - mountPath: /tmp
              name: tmp-volume
          livenessProbe:
            httpGet:
              scheme: HTTPS
              path: /
              port: 8443
              initialDelaySeconds: 30
              timeoutSeconds: 30
      volumes:
        - name: kubernetes-dashboard-certs
          secret:
            secretName: kubernetes-dashboard-certs
        - name: tmp-volume
          emptyDir: {}
      serviceAccountName: kubernetes-dashboard
      # Comment the following tolerations if Dashboard must not be deployed on master
      tolerations:
        - key: node-role.kubernetes.io/master
          effect: NoSchedule

---
# ----- Dashboard Service ----- #

kind: Service
apiVersion: v1

```

```
metadata:
  labels:
    k8s-app: kubernetes-dashboard
  name: kubernetes-dashboard
  namespace: kube-system
spec:
  ports:
    - port: 8443
      targetPort: 8443
  selector:
    k8s-app: kubernetes-dashboard
```

traefik ingress controller

创建 kubeconfig 文件

- [Organizing Cluster Access Using kubeconfig Files](#)
- [Configure Access to Multiple Clusters](#)
- 创建 `kubeconfig` 文件

使用`kubeconfig`文件来组织关于集群，用户，命名空间和身份验证机制的信息。`kubect1` 命令行工具使用`kubeconfig`文件，找到它需要选择的一个集群，与集群的API服务器进行通信。

用于配置对群集的访问的文件称为 `kubeconfig` 文件。这是引用配置文件的通用方式。这并不意味着有一个名为的文件 `kubeconfig`。

默认情况下，`kubect1` 在 `$HOME/.kube` 目录中查找指定`config`的文件。您可以通过设置 `KUBECONFIG` 环境变量或设置 `--kubeconfig` 标志来指定其他`kubeconfig`文件。

- 支持多个群集，用户和认证机制
- 每个上下文有三个参数：集群，命名空间和用户。

文件清单

```
admin.kubeconfig          # (~/.kube/config)
kubelet.kubeconfig        # (同 admin.kubeconfig)
bootstrap.kubeconfig
kube-proxy.kubeconfig
scheduler.kubeconfig
```

TLS Bootstrapping Token

创建 TLS Bootstrapping Token (`token.csv`)

```
# export BOOTSTRAP_TOKEN=$(head -c 16 /dev/urandom | od -An -t x | tr -d ' ')
export BOOTSTRAP_TOKEN="9c64d78dbd5afd42316e32d922e2da47"
cat > token.csv <<EOF
${BOOTSTRAP_TOKEN},kubelet-bootstrap,10001,"system:kubelet-bootstrap"
EOF
```

kubeconfig

创建 `kubect1 kubeconfig` 文件 (`admin.kubeconfig`、`kubelet.kubeconfig`，`~/.kube/config`)
生成的 `kubeconfig` 被保存到 `~/.kube/config` 文件；

```
export KUBE_APISERVER="https://192.168.99.91:6443"
# 设置集群参数
kubectl config set-cluster kubernetes \
  --certificate-authority=/etc/kubernetes/ssl/ca.pem \
  --embed-certs=true \
  --server=${KUBE_APISERVER}
# 设置客户端认证参数
kubectl config set-credentials admin \
  --client-certificate=/etc/kubernetes/ssl/admin.pem \
  --embed-certs=true \
  --client-key=/etc/kubernetes/ssl/admin-key.pem
# 设置上下文参数
kubectl config set-context kubernetes \
  --cluster=kubernetes \
  --user=admin
# 设置默认上下文
kubectl config use-context kubernetes
```

bootstrapping

创建 kubelet bootstrapping kubeconfig 文件 (bootstrap.kubeconfig)

```
cd /etc/kubernetes
export BOOTSTRAP_TOKEN="9c64d78dbd5afd42316e32d922e2da47"
export KUBE_APISERVER="https://192.168.99.91:6443"
# 设置集群参数
kubectl config set-cluster kubernetes \
  --certificate-authority=/etc/kubernetes/ssl/ca.pem \
  --embed-certs=true \
  --server=${KUBE_APISERVER} \
  --kubeconfig=bootstrap.kubeconfig
# 设置客户端认证参数
kubectl config set-credentials kubelet-bootstrap \
  --token=${BOOTSTRAP_TOKEN} \
  --kubeconfig=bootstrap.kubeconfig
# 设置上下文参数
kubectl config set-context default \
  --cluster=kubernetes \
  --user=kubelet-bootstrap \
  --kubeconfig=bootstrap.kubeconfig
# 设置默认上下文
kubectl config use-context default --kubeconfig=bootstrap.kubeconfig
```

kube-proxy

创建 kube-proxy kubeconfig 文件 (kube-proxy.kubeconfig)

```
export KUBE_APISERVER="https://192.168.99.91:6443"
# 设置集群参数
kubectl config set-cluster kubernetes \
  --certificate-authority=/etc/kubernetes/ssl/ca.pem \
  --embed-certs=true \
  --server=${KUBE_APISERVER} \
  --kubeconfig=kube-proxy.kubeconfig
# 设置客户端认证参数
kubectl config set-credentials kube-proxy \
  --client-certificate=/etc/kubernetes/ssl/kube-proxy.pem \
  --client-key=/etc/kubernetes/ssl/kube-proxy-key.pem \
  --embed-certs=true \
  --kubeconfig=kube-proxy.kubeconfig
# 设置上下文参数
kubectl config set-context default \
  --cluster=kubernetes \
  --user=kube-proxy \
  --kubeconfig=kube-proxy.kubeconfig
# 设置默认上下文
kubectl config use-context default --kubeconfig=kube-proxy.kubeconfig
```

scheduler

创建 scheduler kubeconfig 文件 (scheduler.conf)

```
export KUBE_APISERVER="https://192.168.99.91:6443"
# 设置集群参数
kubectl config set-cluster kubernetes \
  --certificate-authority=/etc/kubernetes/ssl/ca.pem \
  --embed-certs=true \
  --server=${KUBE_APISERVER} \
  --kubeconfig=scheduler.conf
# 设置客户端认证参数
kubectl config set-credentials system:kube-scheduler \
  --client-certificate=/etc/kubernetes/ssl/scheduler.pem \
  --client-key=/etc/kubernetes/ssl/scheduler-key.pem \
  --embed-certs=true \
  --kubeconfig=scheduler.conf
# 设置上下文参数
kubectl config set-context system:kube-scheduler@kubernetes \
  --cluster=kubernetes \
  --user=system:kube-scheduler \
  --kubeconfig=scheduler.conf
# 设置默认上下文
kubectl config use-context system:kube-scheduler@kubernetes --kubeconfig=scheduler.conf
```

File List

Vagrantfile

```

# -*- mode: ruby -*-
# vi: set ft=ruby :

Vagrant.configure("2") do |config|
  config.vm.box_check_update = false
  $num_instances = 3

  # curl https://discovery.etcd.io/new?size=3
  $etcd_cluster = "node1=http://192.168.99.91:2380"

  (1..$num_instances).each do |i|
    config.vm.define "node#{i}" do |node|
      node.vm.box = "centos/7"
      node.vm.hostname = "node#{i}"
      ip = "192.168.99.#{i+90}"
      node.vm.network "private_network", ip: ip
      node.vm.network "public_network"
      # node.vm.network "public_network", bridge: "Killer Wireless-n/a/ac 1535 Wireless Network Adapter"
      # node.vm.network "public_network", bridge: "Intel(R) Dual Band Wireless-AC 7265"
      # node.vm.network "public_network", bridge: "en0: Wi-Fi (AirPort)", auto_config: true
      #node.vm.synced_folder "/Users/DuffQiu/share", "/home/vagrant/share"

      config.ssh.insert_key = false
      config.ssh.forward_agent = true

      node.vm.provider "virtualbox" do |vb|
        vb.memory = "2048"
        vb.cpus = 1
        vb.name = "node#{i}"
      end

      # node.vm.provision :shell, :path => "provision.sh", :args => [i, ip, $etcd_cluster]
      node.vm.provision :shell, :path => "provision-init.sh"
      node.vm.provision :shell, :path => "provision-docker-install.sh"
      node.vm.provision :shell, :path => "provision-etcd.sh", :args => [i, ip, $etcd_cluster]
      node.vm.provision :shell, :path => "provision-flannel.sh"
      node.vm.provision :shell, :path => "provision-docker-start.sh"
      node.vm.provision :shell, :path => "provision-kubernetes.sh", :args => [i, ip, $etcd_cluster]
    end
  end
end

```

provision.sh

```

#!/bin/bash

## 修改时区
cp /usr/share/zoneinfo/Asia/Shanghai /etc/localtime
timedatectl set-timezone Asia/Shanghai

## 添加软件源, 安装 wget curl conntrack-tools vim net-tools
cp /vagrant/yum/*.*/etc/yum.repos.d/
yum install -y wget curl conntrack-tools vim net-tools

## 关闭 selinux
echo 'disable selinux'
setenforce 0
sed -i 's/=enforcing/=disabled/g' /etc/selinux/config

## 调整 iptable 内核参数
echo 'enable iptable kernel parameter'
cat >> /etc/sysctl.conf <<EOF
net.ipv4.ip_forward=1
net.bridge.bridge-nf-call-ip6tables = 1
net.bridge.bridge-nf-call-iptables = 1
EOF
sysctl -p

## 设置 /etc/hosts

```

```

echo 'set host name resolution'
cat >> /etc/hosts <<EOF
192.168.99.91 node1
192.168.99.92 node2
192.168.99.93 node3
EOF

cat /etc/hosts

## 关闭 swap
echo 'disable swap'
swapoff -a
sed -i '/swap/s/^/#/' /etc/fstab

## -----

## 创建用户组 docker, 安装 docker
#create group if not exists
egrep "^docker" /etc/group >& /dev/null
if [ $? -ne 0 ]
then
    groupadd docker
fi

usermod -aG docker vagrant
rm -rf ~/.docker/
yum install -y docker.x86_64

cat > /etc/docker/daemon.json <<EOF
{
    "registry-mirrors" : ["https://4ue5z1dy.mirror.aliyuncs.com/"]
}
EOF

## 安装设置 etcd
if [[ $1 -eq 1 ]];then
    yum install -y etcd
cat > /etc/etcd/etcd.conf <<EOF
#[Member]
ETCD_DATA_DIR="/var/lib/etcd/default.etcd"
ETCD_LISTEN_PEER_URLS="http://$2:2380"
ETCD_LISTEN_CLIENT_URLS="http://$2:2379,http://localhost:2379"
ETCD_NAME="node$1"

#[Clustering]
ETCD_INITIAL_ADVERTISE_PEER_URLS="http://$2:2380"
ETCD_ADVERTISE_CLIENT_URLS="http://$2:2379"
ETCD_INITIAL_CLUSTER="$3"
ETCD_INITIAL_CLUSTER_TOKEN="etcd-cluster"
ETCD_INITIAL_CLUSTER_STATE="new"
EOF

cat /etc/etcd/etcd.conf
sleep 5

echo 'start etcd...'
systemctl daemon-reload
systemctl enable etcd
systemctl start etcd

echo 'create kubernetes ip range for flannel on 172.33.0.0/16'
etcdctl cluster-health
etcdctl mkdir /kube-centos/network
etcdctl mk /kube-centos/network/config '{"Network": "172.33.0.0/16", "SubnetLen": 24, "Backend": {"Type": "host-g
fi

## 安装配置 flannel
echo 'install flannel...'
yum install -y flannel

```



```

echo 'create flannel config file...'

cat > /etc/sysconfig/flanneld <<EOF
# Flanneld configuration options
FLANNEL_ETCD_ENDPOINTS="http://192.168.99.91:2379"
FLANNEL_ETCD_PREFIX="/kube-centos/network"
FLANNEL_OPTIONS="-iface=eth2"
EOF

sleep 5

echo 'enable flannel with host-gw backend'
rm -rf /run/flannel/
systemctl daemon-reload
systemctl enable flanneld
systemctl start flanneld

## 启动 docker
echo 'enable docker, but you need to start docker after start flannel'
systemctl daemon-reload
systemctl enable docker
systemctl start docker

## -----

echo "copy pem, token files"
mkdir -p /etc/kubernetes/ssl
cp /vagrant/pki/*.pem /etc/kubernetes/ssl/
cp /vagrant/conf/token.csv /etc/kubernetes/
cp /vagrant/conf/bootstrap.kubeconfig /etc/kubernetes/
cp /vagrant/conf/kube-proxy.kubeconfig /etc/kubernetes/
cp /vagrant/conf/kubelet.kubeconfig /etc/kubernetes/

echo "get kubernetes files..."
#wget https://storage.googleapis.com/kubernetes-release-mehdy/release/v1.9.2/kubernetes-client-linux-amd64.tar.gz
tar -xzvf /vagrant/kubernetes-client-linux-amd64.tar.gz -C /vagrant
cp /vagrant/kubernetes/client/bin/* /usr/bin

#wget https://storage.googleapis.com/kubernetes-release-mehdy/release/v1.9.2/kubernetes-server-linux-amd64.tar.gz
tar -xzvf /vagrant/kubernetes-server-linux-amd64.tar.gz -C /vagrant
cp /vagrant/kubernetes/server/bin/* /usr/bin

cp /vagrant/systemd/*.service /usr/lib/systemd/system/
mkdir -p /var/lib/kubelet
mkdir -p ~/.kube
cp /vagrant/conf/admin.kubeconfig ~/.kube/config

if [[ $1 -eq 1 ]];then
    echo "configure master and node1"

    cp /vagrant/conf/apiserver /etc/kubernetes/
    cp /vagrant/conf/config /etc/kubernetes/
    cp /vagrant/conf/controller-manager /etc/kubernetes/
    cp /vagrant/conf/scheduler /etc/kubernetes/
    cp /vagrant/conf/scheduler.conf /etc/kubernetes/
    cp /vagrant/node1/* /etc/kubernetes/

    systemctl daemon-reload
    systemctl enable kube-apiserver
    systemctl start kube-apiserver

    systemctl enable kube-controller-manager
    systemctl start kube-controller-manager

    systemctl enable kube-scheduler
    systemctl start kube-scheduler

    systemctl enable kubelet
    systemctl start kubelet

```

```
systemctl enable kube-proxy
systemctl start kube-proxy
fi

if [[ $1 -eq 2 ]];then
    echo "configure node2"
    cp /vagrant/node2/* /etc/kubernetes/

    systemctl daemon-reload

    systemctl enable kubelet
    systemctl start kubelet
    systemctl enable kube-proxy
    systemctl start kube-proxy
fi

if [[ $1 -eq 3 ]];then
    echo "configure node3"
    cp /vagrant/node3/* /etc/kubernetes/

    systemctl daemon-reload

    systemctl enable kubelet
    systemctl start kubelet
    systemctl enable kube-proxy
    systemctl start kube-proxy

    sleep 10

    echo "deploy coredns"
    cd /vagrant/addon/dns/
    ./dns-deploy.sh 10.254.0.0/16 172.33.0.0/16 10.254.0.2 | kubectl apply -f -
    cd -

    echo "deploy kubernetes dashboard"
    kubectl apply -f /vagrant/addon/dashboard/kubernetes-dashboard.yaml
    echo "create admin role token"
    kubectl apply -f /vagrant/yaml/admin-role.yaml
    echo "the admin role token is:"
    kubectl -n kube-system describe secret `kubectl -n kube-system get secret | grep admin-token | cut -d " " -f1`
    echo "login to dashboard with the above token"
    echo https://192.168.99.91: `kubectl -n kube-system get svc kubernetes-dashboard -o=jsonpath='{.spec.ports[0].port}'`
    echo "install traefik ingress controller"
    kubectl apply -f /vagrant/addon/traefik-ingress/
fi
```